



The University of Manchester

Running Parallel Bytecode Interpreters on Heterogeneous Hardware

Juan Fumero, Athanasios Stratikopoulos, Christos Kotselidis

MoreVMs 2020



<juanfumero@acm.org>

University of Manchester

Outline

- Motivation
- Parallel Bytecode Interpreter
 - Initial implementation
 - Parallel Implementation
 - Multiple-heap configuration
- Initial Results
- Takeaways

Heterogeneous Hardware is Everywhere



Motivation

RQ: Heterogeneous systems are everywhere.

- a) Can we run an Interpreter on multiple Heterogeneous Devices?**
- b) Can we increase performance?**

- Accelerate components of the actual VMs
 - Garbage Collection (GC)

[US Patent 2010 0082930 A1] [GPU Assisted Garbage Collection](#) (AMD)

[ISMM'12] Offloading Garbage Collection on GPUs

[CASES'16] Generational GCs on Integrated GPUs (FastCollect)

Motivation

RQ: Heterogeneous systems are everywhere.

- a) Can we run an Interpreter on multiple Heterogeneous Devices?
- b) Can we increase performance?

- Accelerate components of the actual VMs
 - GC?
 - **Bytecode interpreters**

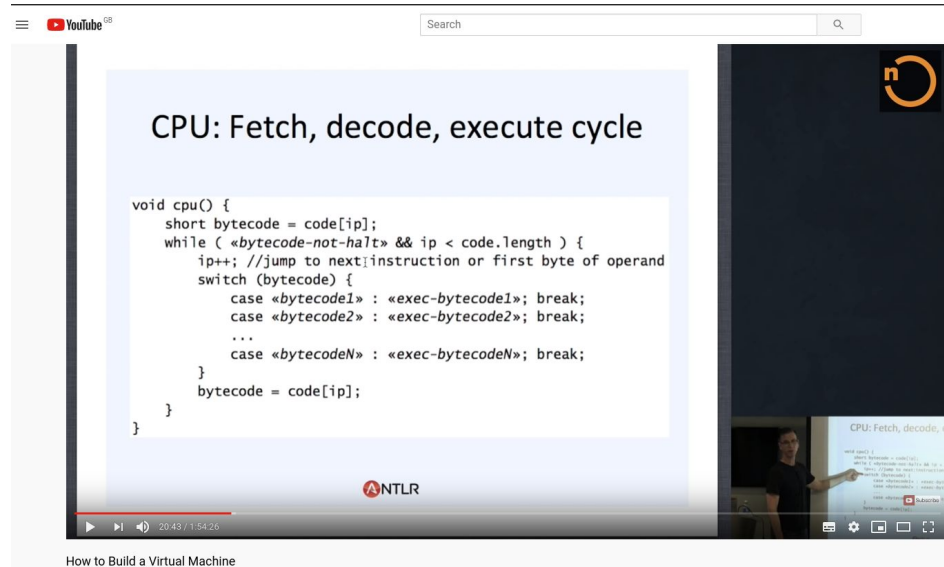


In this work, we focus on the
bytecode interpreter

[OOPSLA'19] CUDA Single Thread Interpreters

Bytecode Interpreter

Subset of Java bytecodes → Toy example, but simple and powerful enough to start computing some workloads



The image shows a YouTube video player interface. The video content is a presentation slide titled "CPU: Fetch, decode, execute cycle". The slide contains C code for a CPU function. The code is as follows:

```
void cpu() {  
    short bytecode = code[ip];  
    while ( «bytecode-not-halt» && ip < code.length ) {  
        ip++; //jump to next instruction or first byte of operand  
        switch (bytecode) {  
            case «bytecode1» : «exec-bytecode1»; break;  
            case «bytecode2» : «exec-bytecode2»; break;  
            ...  
            case «bytecodeN» : «exec-bytecodeN»; break;  
        }  
        bytecode = code[ip];  
    }  
}
```

The video player shows a person pointing at the slide. The video title is "How to Build a Virtual Machine". The video player controls show a progress bar at 20:43 / 1:54:26.

Example BC Interpreter by Professor Terence Parr

Simple stack-based machine

We extend it with more bytecodes and port to OpenCL

<https://www.youtube.com/watch?v=OjaAToVkoTw>

Bytecode Interpreter

- Arithmetic Operations: IDIV, IADD, IMUL, ISUB
- Bitwise: RSHIFT, LSHIFT
- Comparisons: ILT, IEQ
- Memory: STORE, LOAD, GSTORE_INDEXED, GLOAD_INDEXED,
- Control Flow: BR, BRT, BRF, HALT, RET, CALL
- Interpreter Control: POP, DUP, ICONST1, ICONST <n>
- Auxiliar: PRINT

Bytecode Interpreter

```
// Expressing vector multiplication
```

```
ICONST, 0,  
DUP,  
ICONST, SIZE,  
IEQ,  
BRT, 23,  
DUP,  
DUP,  
GLOAD_INDEXED, SIZE,  
LOAD, 1,  
GLOAD_INDEXED, SIZE * 2,  
IMUL,  
GSTORE_INDEXED, BASE,  
ICONST1,  
IADD,  
BR, 2,  
POP,  
HALT
```


Bytecode Interpreter

```
// Expressing vector multiplication
```

```
ICONST, 0,  
DUP,  
ICONST, SIZE,  
IEQ,  
BRT, 23,  
DUP,  
DUP,  
GLOAD_INDEXED, SIZE,  
LOAD, 1,  
GLOAD_INDEXED, SIZE * 2,  
IMUL,  
GSTORE_INDEXED, BASE,  
ICONST1,  
IADD,  
BR, 2,  
POP,  
HALT
```

HEAP



0

Bytecode Interpreter

```
// Expressing vector multiplication
```

```
ICONST, 0,
```

```
DUP,
```

```
ICONST, SIZE,
```

```
IEQ,
```

```
BRT, 23,
```

```
DUP,
```

```
DUP,
```

```
GLOAD_INDEXED, SIZE,
```

```
LOAD, 1,
```

```
GLOAD_INDEXED, SIZE * 2,
```

```
IMUL,
```

```
GSTORE_INDEXED, BASE,
```

```
ICONST1,
```

```
IADD,
```

```
BR, 2,
```

```
POP,
```


```
HALT
```

HEAP

A


B

C



0

0



Bytecode Interpreter

// Expressing vector multiplication

```

ICONST, 0,
DUP,
ICONST, SIZE,
IEQ,
BRT, 23,
DUP,
DUP,
GLOAD_INDEXED, SIZE,
LOAD, 1,
GLOAD_INDEXED, SIZE * 2,
IMUL,
GSTORE_INDEXED, BASE,
ICONST1,
IADD,
BR, 2,
POP,
HALT
    
```

HEAP



SIZE
0
0

Bytecode Interpreter

```
// Expressing vector multiplication
```

```
ICONST, 0,  
DUP,  
ICONST, SIZE,  
IEQ,  
BRT, 23,  
DUP,  
DUP,  
GLOAD_INDEXED, SIZE,  
LOAD, 1,  
GLOAD_INDEXED, SIZE * 2,  
IMUL,  
GSTORE_INDEXED, BASE,  
ICONST1,  
IADD,  
BR, 2,  
POP,  
HALT
```

HEAP



SIZE
0
0

SIZE == 0? GOTO 23: next

Bytecode Interpreter

```
// Expressing vector multiplication
```

```
ICONST, 0,  
DUP,  
ICONST, SIZE,  
IEQ,  
BRT, 23,  
DUP,  
DUP,  
GLOAD_INDEXED, SIZE,  
LOAD, 1,  
GLOAD_INDEXED, SIZE * 2,  
IMUL,  
GSTORE_INDEXED, BASE,  
ICONST1,  
IADD,  
BR, 2,  
POP,  
HALT
```

HEAP



0


Bytecode Interpreter

// Expressing vector multiplication

```

ICONST, 0,
DUP,
ICONST, SIZE,
IEQ,
BRT, 23,
DUP,
DUP,
GLOAD_INDEXED, SIZE,
LOAD, 1,
GLOAD_INDEXED, SIZE * 2,
IMUL,
GSTORE_INDEXED, BASE,
ICONST1,
IADD,
BR, 2,
POP,
HALT
    
```

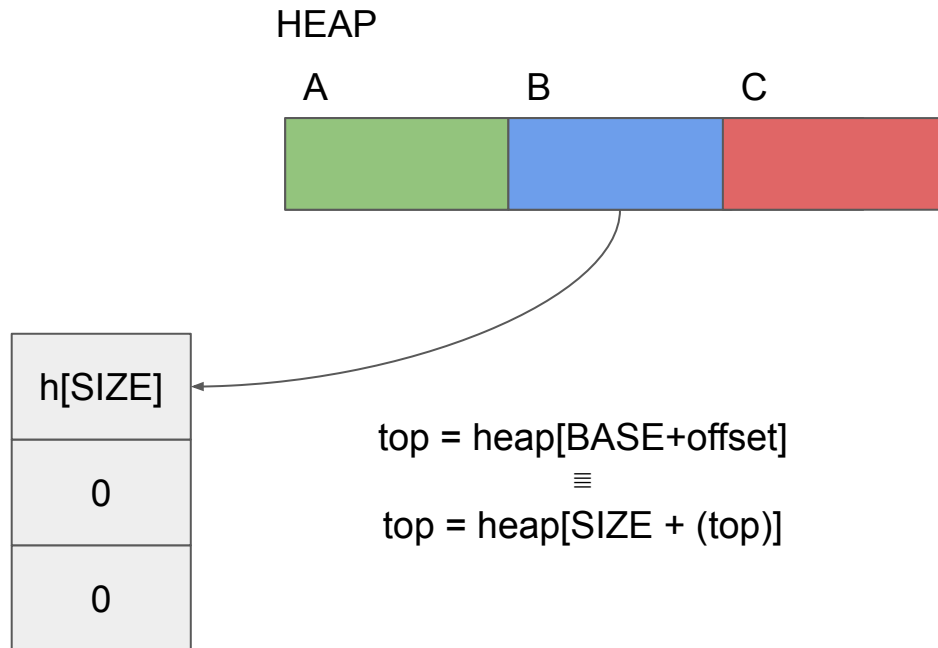
HEAP



Bytecode Interpreter

```
// Expressing vector multiplication
```


```
ICONST, 0,  
DUP,  
ICONST, SIZE,  
IEQ,  
BRT, 23,  
DUP,  
DUP,  
GLOAD_INDEXED, SIZE,  
LOAD, 1,  
GLOAD_INDEXED, SIZE * 2,  
IMUL,  
GSTORE_INDEXED, BASE,  
ICONST1,  
IADD,  
BR, 2,  
POP,  
HALT
```



Bytecode Interpreter

```
// Expressing vector multiplication
```

```
ICONST, 0,  
DUP,  
ICONST, SIZE,  
IEQ,  
BRT, 23,  
DUP,  
DUP,  
GLOAD_INDEXED, SIZE,  
LOAD, 1,  
GLOAD_INDEXED, SIZE * 2,  
IMUL,  
GSTORE_INDEXED, BASE,  
ICONST1,  
IADD,  
BR, 2,  
POP,  
HALT
```



HEAP



Bytecode Interpreter

// Expressing vector multiplication

```

ICONST, 0,
DUP,
ICONST, SIZE,
IEQ,
BRT, 23,
DUP,
DUP,
GLOAD_INDEXED, SIZE,
LOAD, 1,
GLOAD_INDEXED, SIZE * 2,
IMUL,
GSTORE_INDEXED, BASE,
ICONST1,
IADD,
BR, 2,
POP,
HALT
    
```

HEAP



$\text{top} = \text{heap}[\text{BASE} + \text{offset}]$

\equiv

$\text{top} = \text{heap}[\text{SIZE} + (\text{top})]$


Bytecode Interpreter

// Expressing vector multiplication

```

ICONST, 0,
DUP,
ICONST, SIZE,
IEQ,
BRT, 23,
DUP,
DUP,
GLOAD_INDEXED, SIZE,
LOAD, 1,
GLOAD_INDEXED, SIZE * 2,
IMUL,
GSTORE_INDEXED, BASE,
ICONST1,
IADD,
BR, 2,
POP,
HALT
    
```

HEAP




Bytecode Interpreter

```
// Expressing vector multiplication
```

```
ICONST, 0,  
DUP,  
ICONST, SIZE,  
IEQ,  
BRT, 23,  
DUP,  
DUP,  
GLOAD_INDEXED, SIZE,  
LOAD, 1,  
GLOAD_INDEXED, SIZE * 2,  
IMUL,  
GSTORE_INDEXED, BASE,  
ICONST1,  
IADD,  
BR, 2,  
POP,  
HALT
```

HEAP



$\text{heap}[\text{BASE} + \text{offset}] = \text{top}$
 $=$
 $\text{heap}[\text{BASE} + (\text{top} - 1)] = \text{top}$

0


Bytecode Interpreter

// Expressing vector multiplication

```

ICONST, 0,
DUP,
ICONST, SIZE,
IEQ,
BRT, 23,
DUP,
DUP,
GLOAD_INDEXED, SIZE,
LOAD, 1,
GLOAD_INDEXED, SIZE * 2,
IMUL,
GSTORE_INDEXED, BASE,
ICONST1,
IADD,
BR, 2,
POP,
HALT
    
```

HEAP



Bytecode Interpreter

```
// Expressing vector multiplication
```

```
ICONST, 0,  
DUP,  
ICONST, SIZE,  
IEQ,  
BRT, 23,  
DUP,  
DUP,  
GLOAD_INDEXED, SIZE,  
LOAD, 1,  
GLOAD_INDEXED, SIZE * 2,  
IMUL,  
GSTORE_INDEXED, BASE,  
ICONST1,  
IADD,  
BR, 2,  
POP,  
HALT
```

HEAP



Bytecode Interpreter

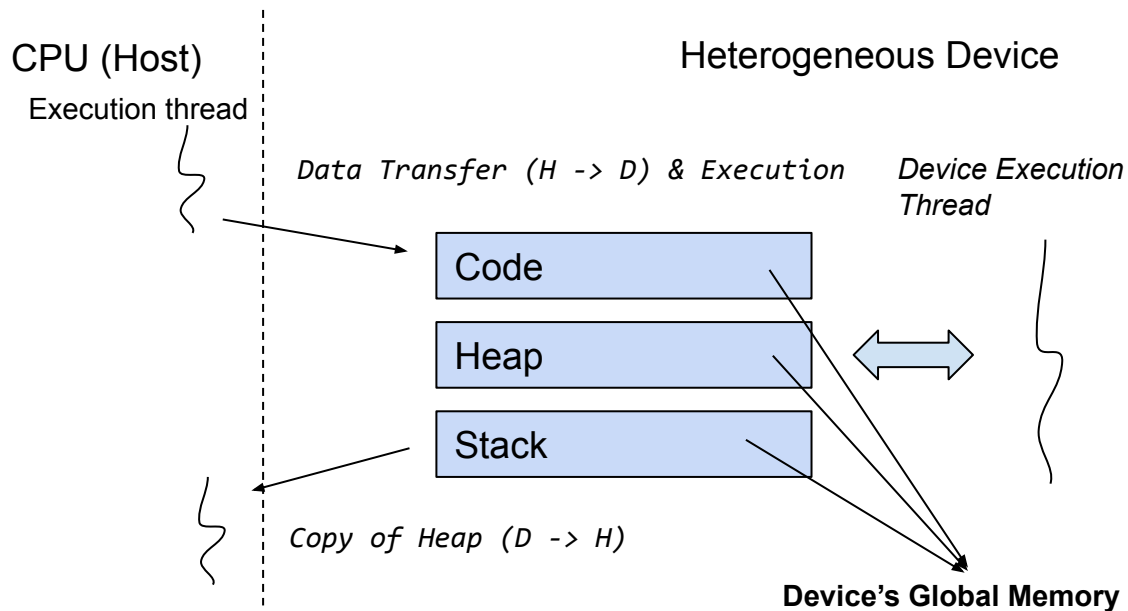
```
// Expressing vector multiplication
```

```
ICONST, 0,  
DUP,  
ICONST, SIZE,  
IEQ,  
BRT, 23,  
DUP,  
DUP,  
GLOAD_INDEXED, SIZE,  
LOAD, 1,  
GLOAD_INDEXED, SIZE * 2,  
IMUL,  
GSTORE_INDEXED, BASE,  
ICONST1,  
IADD,  
BR, 2,  
POP,  
HALT
```

HEAP



OpenCL Bytecode Interpreter



OpenCL Bytecode Interpreter


```
__attribute__((num_compute_units(1)))  
__attribute__((reqd_work_group_size(1,1,1)))  
__kernel void interpreter(global int* code, global int* stack,  
                          global int* data, global char* buffer,  
                          const int codeSize, int ip, int fp, int sp, int trace) {  
    while (ip < codeSize) {  
        int opcode = code[ip];  
        ip++;  
        switch (opcode) {  
            case BC1: ... break;  
            case BC2: ... break;  
            case BC3: ... break;  
            case BC4: ... break;  
            case BC5: ... break;  
            case BC6: ... break;  
        }  
    }  
}
```



It will be open-source soon!

Initial Results

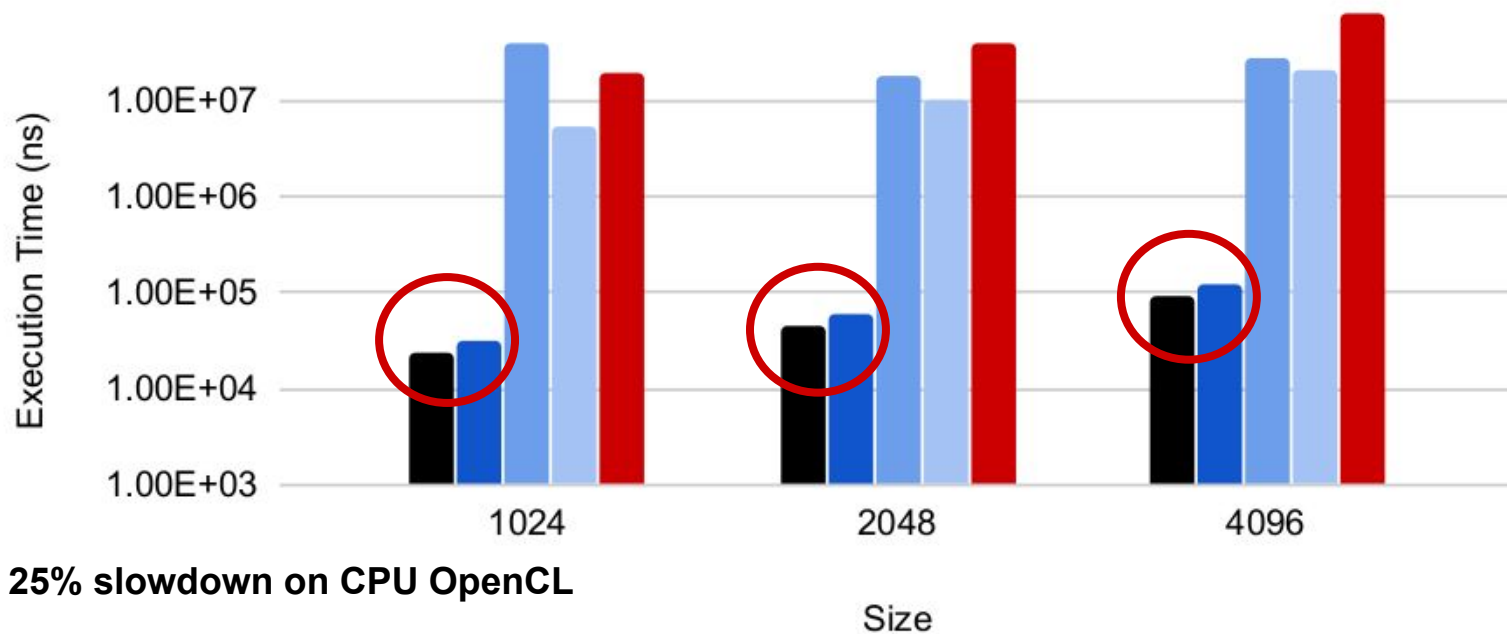
■ C++ on Intel CPU ■ Intel CPU OpenCL ■ NVIDIA GP100 Single Thread
■ Intel HD Graphics Single Thread ■ Xilinx FPGA Single Thread



The lower, the better

Initial Results

■ C++ on Intel CPU ■ Intel CPU OpenCL ■ NVIDIA GP100 Single Thread
■ Intel HD Graphics Single Thread ■ Xilinx FPGA Single Thread



Only 25% slowdown on CPU OpenCL


Can we do better?

Optimising for Heterogeneous Hardware

- A) Parallel Interpreter through the introduction of Thread-Identifier
- B) Memory regions (tier memory)

A) Thread-ID in the Interpreter

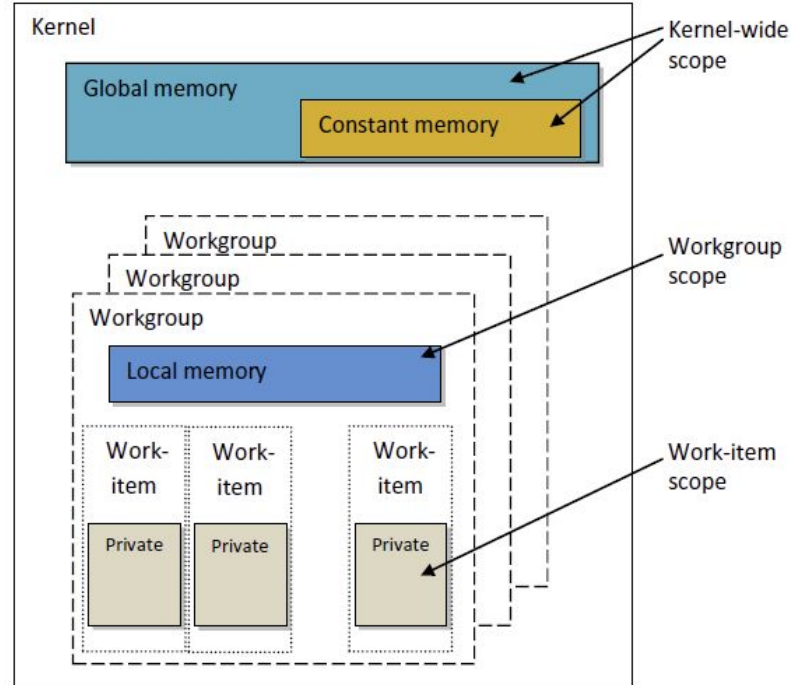
```
case THREAD_ID:  
    value = get_local_id(0);  
    stack[++sp] = value;  
    break;
```



ND-Range

THREAD_ID

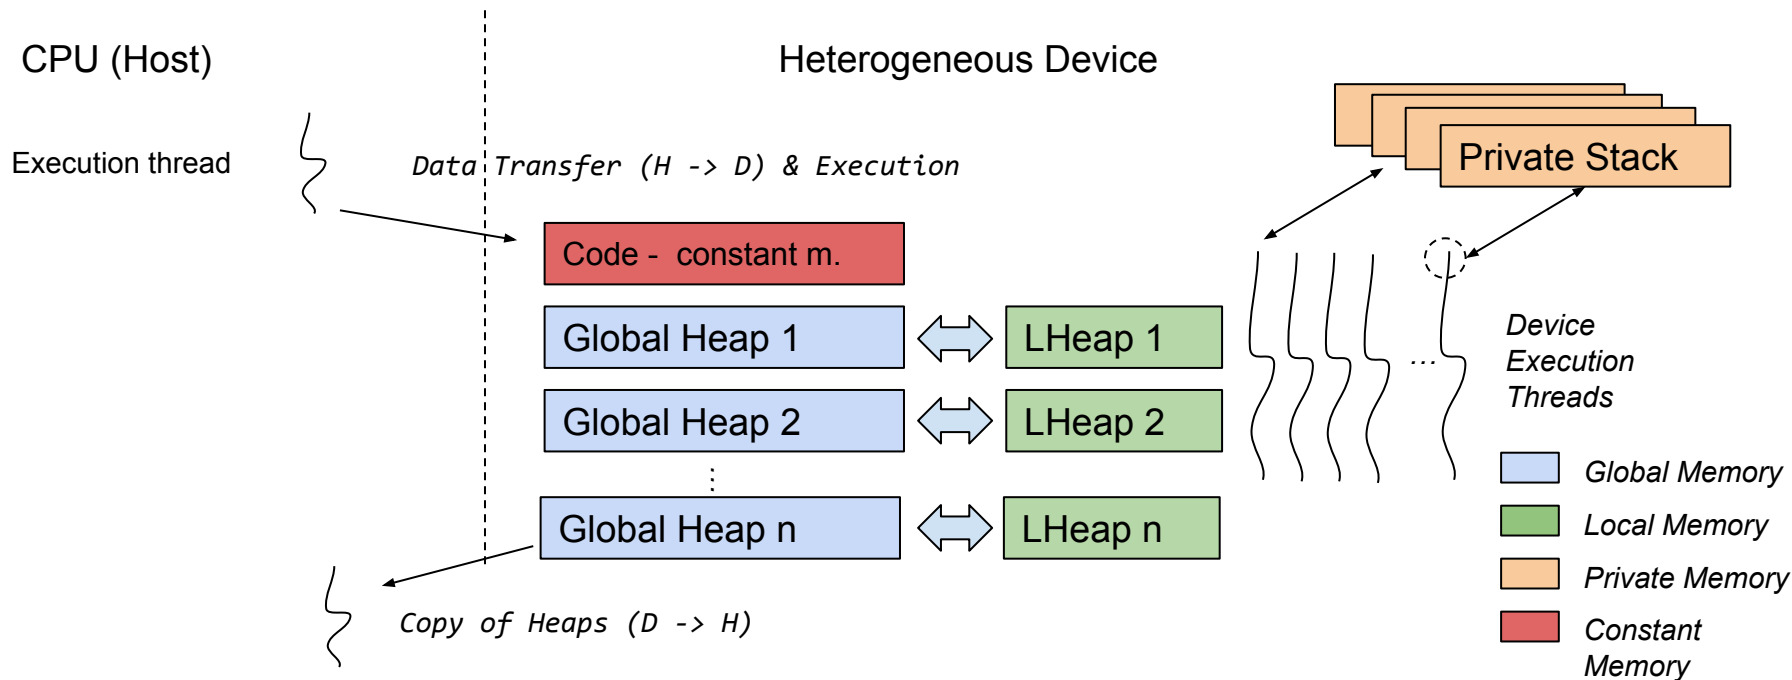
Understanding OpenCL MM



- Exploit Memory Regions

Source: <https://www.mql5.com/en/articles/407>

B) Memory Regions for out BC-Interpreter



OpenCL Bytecode Interpreter (II)

// Expressing vector multiplication

```
THREAD_ID,  
DUP,  
PARALLEL_GLOAD_INDEXED, 0,  
THREAD_ID,  
PARALLEL_GLOAD_INDEXED, 1,  
IMUL,  
PARALLEL_GSTORE_INDEXED, 2,  
HALT
```

- One stack per thread
- Code in constant memory
- Multiple Global heaps
- Stack in private memory


OpenCL Bytecode Interpreter (II)

// Expressing vector addition

```
THREAD_ID,  
DUP,  
PARALLEL_GLOAD_INDEXED, 0,  
THREAD_ID,  
PARALLEL_GLOAD_INDEXED, 1,  
IMUL,  
PARALLEL_GSTORE_INDEXED, 2,  
HALT
```



Heaps



OpenCL Bytecode Interpreter (II)


// Expressing vector addition

```
THREAD_ID,  
DUP,  
PARALLEL_GLOAD_INDEXED, 0,  
THREAD_ID,  
PARALLEL_GLOAD_INDEXED, 1,  
IMUL,  
PARALLEL_GSTORE_INDEXED, 2,  
HALT
```




120

Heaps




OpenCL Bytecode Interpreter (II)

```
// Expressing vector addition  
THREAD_ID,  
DUP,  
PARALLEL_GLOAD_INDEXED, 0,  
THREAD_ID,  
PARALLEL_GLOAD_INDEXED, 1,  
IMUL,  
PARALLEL_GSTORE_INDEXED, 2,  
HALT
```



Heaps



OpenCL Bytecode Interpreter (II)

```
// Expressing vector addition
```

```
THREAD_ID,
```

```
DUP,
```

```
PARALLEL_GLOAD_INDEXED, 0,
```


```
THREAD_ID,
```

```
PARALLEL_GLOAD_INDEXED, 1,
```

```
IMUL,
```

```
PARALLEL_GSTORE_INDEXED, 2,
```

```
HALT
```



OpenCL Bytecode Interpreter (II)

```
// Expressing vector addition
```

```
THREAD_ID,
```

```
DUP,
```

```
PARALLEL_GLOAD_INDEXED, 0,
```


```
THREAD_ID,
```

```
PARALLEL_GLOAD_INDEXED, 1,
```


```
IMUL,
```

```
PARALLEL_GSTORE_INDEXED, 2,
```

```
HALT
```



Heaps



OpenCL Bytecode Interpreter (II)

// Expressing vector addition

THREAD_ID,

DUP,

PARALLEL_GLOAD_INDEXED, 0,


THREAD_ID,

PARALLEL_GLOAD_INDEXED, 1,

IMUL,

PARALLEL_GSTORE_INDEXED, 2,


HALT




OpenCL Bytecode Interpreter (II)

// Expressing vector addition

```
THREAD_ID,  
DUP,  
PARALLEL_GLOAD_INDEXED, 0,  
THREAD_ID,  
PARALLEL_GLOAD_INDEXED, 1,  
IMUL,  
PARALLEL_GSTORE_INDEXED, 2,  
HALT
```




Heaps



OpenCL Bytecode Interpreter (II)


```
// Expressing vector addition  
THREAD_ID,  
DUP,  
PARALLEL_GLOAD_INDEXED, 0,  
THREAD_ID,  
PARALLEL_GLOAD_INDEXED, 1,  
IMUL,  
PARALLEL_GSTORE_INDEXED, 2,  
HALT
```




OpenCL Bytecode Interpreter (II)

// Expressing vector addition

```
THREAD_ID,  
DUP,  
PARALLEL_GLOAD_INDEXED, 0,  
THREAD_ID,  
PARALLEL_GLOAD_INDEXED, 1,  
IMUL,  
PARALLEL_GSTORE_INDEXED, 2,  
HALT
```



Initial Results



The higher, the better


Size

Initial Results

Compared to ARM 1GHz:

~151x on Intel HD Graphics

~214x on NVIDIA



GPUs:

Compared to Intel i7:

- 11x on Intel HD Graphics
- 17x on NVIDIA GP100

FPGAs:


Slowdown on FPGAs:

- Not enough space
- Hard to tune

How this can be useful?

- 1) Main CPUs with “space” for hardware specialization
 - a) FPGAs inside the CPU
 - b) Space for custom instructions (e.g., ARM Custom Instructions - ARM Cortex M33 -)
- 2) If workloads follow SIMT patterns and/or pipeline computation → Execution on parallel bytecode interpreters can be feasible
- 3) Existing VMs
 - a) E.g., TornadoVM
 - i) When the VM is compiling the code (JIT code to FPGA) → Use the parallel bytecode interpreter
 - ii) Multiple heaps on heterogeneous hardware can be extremely useful (local, constant, etc).

Takeaways

- Heterogeneous hardware is here to stay
- Managed runtime systems could potentially be accelerated using heterogeneous hardware
 - Garbage Collection [ISMM'12], [CASES'16]
 - Bytecode interpreters  This work
 - Other components?
- Promising speedups even for simple examples
 - Parallel Bytecode interpreter
 - Multiple memory regions (full tier memory on the target device)

Future Work

- Include more benchmarks
- Include more analysis for each memory region
- Comparisons against mainstream programming languages
- Use other standards such as SYCL C++ from Khronos Group
 - Intel DPC++
 - Codeplay ComputeCPP

Thank you so much for your attention

This work is partially supported by the EU Horizon 2020 E2Data 780245.



Contact: Juan Fumero <juanfumero@acm.org>

Website: <https://jjfumero.github.io>